



EXPOUND

EXPOUND LABORATORIES™

Finality Assurance™: Can Your Enterprise Compute What It May Rely On?

AI made production abundant. The harder question is which specific result a specific party may rely upon, for a specific purpose, right now.

Kenneth Lawrence Cheney

Expound Laboratories™ · expound.com

Version 1.0 · August 26, 2026

© 2026 Kenneth Lawrence Cheney. All rights reserved.

1. Your system says **APPROVED**. What does that establish?

Every enterprise runs on a small vocabulary of finished. **APPROVED. VERIFIED. COMPLETE. CLOSED. RELEASED. GREEN.** Pick one and ask what it establishes. A record proves only that somebody, or some workflow, put that word in that field. Not that the work was authorized. Not that the intended thing happened. Not that the required evidence exists and applies here. Not that this party may act. Not that the basis holds.

Six things it might establish. It carries one, about the writer.

The word also does not change when the world does. A policy is superseded, an approval revoked, an instrument drifts out of calibration, and the field still reads **APPROVED**, byte for byte identical. It is a memory of a past determination asked to behave like a present measurement.

The problem is not that enterprise systems lack status. It is that status is being asked to carry reliance.

2. Finality Assurance, and the one question to remember

Work is consequential when something downstream acts on it and acting wrongly costs something: money moves, a patient is treated, a batch ships. It qualifies here only where the conditions can be written down and the evidence discharging them named, a test per decision, not per industry. I use **Finality Assurance™** to name my own systems and methods for computing whether such work may be relied upon.

Generation answers **what was produced**. Evaluation answers **what the evidence establishes**. Nothing answers **under what conditions anyone may act on it**, the one an enterprise acts on:

What may a named party rely upon, about a named result, for a named purpose, now?

Named party, because a supplier may rely where a regulator may not. **Named result**, because reliance attaches to an instance, not a reputation. **Named purpose**, because evidence enough to file may not be enough to pay. **Now**, because yesterday's answer may be history.

3. Why the old operating model cannot answer it

Three situations an enterprise must tell apart.

Permitted, but it never happened. A \$40,000 refund is approved. No money leaves the account, and the file still reads **APPROVED**.

It happened, but nobody was permitted to do it. The money leaves. The authority behind it does not.

It happened correctly, and this party may still not rely on it. The evidence was gathered to close the books, and the person asking is the auditor.

A single DONE cannot hold those apart, and neither can a richer vocabulary. The difficulty is not the length of the word list but that the field can be written without regard to the facts.

When genuinely different facts are collapsed into one writable value, the distinctions are not preserved in that value and cannot be recovered from it.

Not "hard to read back." Not present. What an organization recovers, it recovers elsewhere, late, expensively, and weaker than capture at the time.

4. What it means to compute reliance




Figure 1 From declared status to computed reliance

Computed reliance is the answer worked out fresh rather than stored. It changes six things.

Decide what must be true before the work can be relied upon. The requirements are fixed in advance, as a list, for that kind of decision. Requirements invented afterward describe what happened.

Gather evidence against those requirements. Evidence is typed by what it can actually establish: that something was authorized, that it occurred, that its basis is current, and the conditions under which reliance holds. A passing test establishes that a test passed, not that it discharged the obligation. What may be said about a result and what may be acted on are two permissions.

Separate the producer's account from what occurred. "I did it" and "it happened" are different propositions. A system that grades its own work is reporting on itself, which can be truthful and still not be what a consequential decision needs. Tooling a team writes and runs to grade its own work is that, one layer down.

Compute the answer for a particular party and purpose. There is no universal DONE licensing everyone. The output answers one question rather than labels a record.

Make the answer unwritable. People change the evidence and the governing inputs, but nobody types the verdict. Where a conclusion can be written without regard to the facts it represents, the write establishes only that somebody wrote it. Where a write path is genuinely constrained by those facts, that is a different design, and an open one.

The obvious question is who writes the requirement list, since the same objection applies to it. The list is authored in advance, before any result exists, under a separated authority, versioned, and recorded as part of what the decision rests on. Human authorship is not eliminated, only moved where it can be seen and argued with.

Recompute when the basis changes. Historical validity does not silently remain present authority. When a dependency fails, reliance is withdrawn rather than outliving its support, and the parties who relied are told. Withdrawn is not stopped: the answer is graded, so it can say we no longer know this to the standard this decision required, which is not we know this is unsafe. The honest bound is observability. The answer is only as current as the declared observation paths, and a change no path can see stays outside what it claims to know. Evidence nobody gathered does not enter it.

Two objects follow. The **Finality Account** is the case file for one decision. The **Finality Verdict** is the answer computed from it for one question, and means nothing away from it.

One rule is worth carrying away, because it is where this parts company with scoring. It is a choice this architecture makes and defends rather than something proved, and it governs lists where every condition is mandatory.

Bad news wins. A single unmet mandatory condition cannot be averaged away by strength somewhere else. No count, quorum, average, or confidence value raises a result that any admitted evaluation computed as non-passing.

Admitted evaluations are those declared in advance, on that list, as competent to answer the obligation. One nobody declared cannot vote, and one declared cannot be dropped when its answer is unwelcome. Nor can the list be shortened by whoever asks; a list requiring nothing is refused rather than passed. A second opinion may narrow, cap or refuse an answer, never raise it.

5. Four industries. The same missing computation.




Figure 2 Four industries. The same missing computation.

Some of this will sound familiar. Access control decides who may act, never whether the thing happened. Provenance records how something was made, not who may rely on it. Revocation cancels a credential, not reliance on a result. Segregation of duties keeps the producer from grading the work. Tooling governing AI behavior and risk posture governs the producer, not acceptance of what the producer made. This uses them all. In a review of published product documentation bounded to August 2026, none computes the object in question: a present, party- and purpose-scoped decision about whether one already-produced result may be relied upon. That finding is dated and open to re-testing.

These four are illustrations, not results.

6. Why AI makes this urgent

AI did not create this problem. It removed the condition that let enterprises live with it. While work was produced at human pace it could be checked at that pace, and the gap between "a

system says this is finished" and "someone established that it is" was absorbed by people working by hand. Machine production breaks that proportion: output stops tracking headcount while checking still does. Queueing and sampling change how the shortfall is distributed rather than removing it. That is an inferred mechanism, not a measured cause.

What does not scale is the judgment: deciding what must be true, and what counts as evidence that it is. That work is per decision type, not per result, so it is amortized rather than paid each time. Amortized is not cheap. Getting those conditions right for one regulated domain is the dominant cost, and the part least solved.

Speed is only half of it. When software acts on its own, no human decision interval exists unless the system opens one. Approval gates and pre-execution checks open exactly that interval, and there the reliance question goes unanswered.

There is a second reason:

Model accuracy is a claim about a population. Reliance is a claim about this result, and it is not an authority the population figure can grant.

A better model is wrong less often, and nobody can say which cases remain. Methods that estimate risk for one output are real work. What none produces is a permission.

A 2026 review of 257 studies of agentic systems found that keeping the supporting evidence valid and current after they ship remains under-developed. It was written with no knowledge of this work.

7. What changes for an enterprise

These follow if the architecture holds and is adopted. None is a measured observation about any organization.

Delegation and risk. The question could stop being how much AI can be deployed and become which consequential work can be handed to machines because reliance on the result can be established mechanically. It runs the other way too: stale, wrongly scoped or unauthorized work would stop acquiring authority because a workflow moved past it.

Audit and compliance. Much assurance work is retrospective assembly, because operational records are optimized for running the business rather than supporting reliance on one result. When the decision record is produced alongside the work, the question changes from what can be pieced together to what was required, what discharged it, and who relied on it.

Architecture. On consequential decision paths, tests, scans, traces, provenance records and governance tooling each establish something real. **They would become evidence producers rather than authorities over finality.** What a deployment may claim depends on where it sits: one that only reads what others produced can report that something was unsupported; only one placed where it can stop the work may say it prevented anything.

Economics. Finality Assurance computes whether reliance is warranted. Accepted Work per Dollar sets the weighted work an organization was entitled to count against the full cost of everything it attempted, including work refused, retried or abandoned. Correct work refused for want of evidence is a real cost, counted rather than netted out. Without a computed acceptance test there is no accepted set, so the loop needs both.

8. What Finality Assurance™ does not mean

Finality Assurance is not a machine for establishing truth. It does not replace human judgment, domain safety engineering or regulatory authority. Safety-critical control is a different category and is not claimed: it is blocked by hazard analysis, timing determinism, independent protective barriers and regulatory approval, none of which this supplies or reduces. It is not a confidence score, nor an argument that every task should be governed. A producer's self-assessment does not establish that a third party may rely.

Closing the write path to the verdict moves trust rather than removing it. It moves to the requirement list, the definitions behind it, the adapters that turn other systems' output into evidence, and the records of who holds authority. An unnamed trust surface cannot be audited or separated from the party it protects.

The claim is narrower: reliance on consequential work can be computed from declared requirements and typed evidence, under a named party and purpose, sound relative to a declared model, and recomputed when the basis moves.

9. The first test an enterprise can run

Take one consequential decision your organization repeats. What must be true before anyone may rely on it? What evidence establishes each? Who decides the requirements? Who may rely, and for what? What could invalidate the answer tomorrow, and who would know?

Most organizations answer some in conversation, few without a person vouching for the answer. That gap is the subject.

The technical paper states separately what has been implemented, what has been measured, and what remains unclaimed; prototype evidence there is not evidence of production readiness. The first move costs nothing: find out whether your enterprise declares answers it could compute.

Before an organization delegates more consequential work to machines, it should know how it will compute what may be relied upon. That is an architecture question, not an AI feature.

Sources

External sources referenced in this article, retrieved on the access date shown. Third parties, standards and publications named here are referred to descriptively and nominatively. Nothing implies endorsement, partnership, implementation, adoption or certification by any of them, and third-party names are the property of their respective owners.

1. **Mirto, F. O., D'Agati, L., Tricomi, G., Silvestri, S., Longo, F., Puliafito, A., Merlino, G.** *Beyond Component Testing: Validating Agentic AI Systems*. Synthesis of 257 studies; behavioral evaluation comparatively mature, temporal validity and runtime evidence maintenance under-developed. arXiv:2607.29405, July 31, 2026. <https://arxiv.org/abs/2607.29405> · accessed 2026-08-26

2. **PCAOB, AS 1105, Audit Evidence.** Sufficiency and appropriateness; more of the same poor-quality evidence cannot compensate for its quality; independent knowledgeable sources are more reliable than internal ones. <https://pcaobus.org/oversight/standards/auditing-standards/details/AS1105> · accessed 2026-08-26

On the metric. The numerator of Accepted Work per Dollar™ depends on a governed value weight, which makes the weights a governance surface: whoever sets them can move the figure without changing any operational result, so the mechanism being evaluated must not set them. Two figures are comparable only when the same acceptance test, the same weights and the same cost model stand behind both. The [technical paper](#) states the discipline that goes with the metric.

Availability. The Finality Assurance family of standards is a proprietary specification family and not a consensus or accredited standard. It is available for licensing under commercial terms. Separately, for the governed routing and constrained-learning layer named in the technical paper and not the standards family above, a public open-source release is planned, with final license terms subject to legal and release review. Treat statements of licensing, availability and timing here as forward-looking rather than committed. Nothing in this article is an offer, a commitment, or a guarantee of terms.

On publication and rights. The concepts, questions, examples and diagnostic vocabulary in this article are published for discussion, testing, teaching and critique. Nothing in this publication grants a license under any patent or patent application, copyright, trademark or other intellectual-

property right, and it grants no implementation, certification or commercial license of any kind. These rights are reserved. The architecture is the subject of United States patent applications filed before the date of this publication, and also includes separately retained trade secrets, copyrighted materials and marks. No patent has issued, no examination has concluded, and nothing here asserts that any claim is valid, that any party infringes, or that any party is free to operate. Certain implementation mechanisms remain confidential and are not disclosed here. Any future open-source release will be governed by the license accompanying that release.