

Who Can Afford Friction? How AI Reprices Economic Power

Friction, Change Capacity, and Price

An economic point of view on how artificial intelligence changes the cost, incidence, and strategic value of business friction: the two-sided mechanism, a balance sheet and five analytical tools, organizational change capacity, and price as the third lens.

AI is not making the economy frictionless. It is changing who can afford friction.

by Kenneth L. Cheney

<https://expound.com> • info@expound.com

Public Article • August 2026

Evidence cutoff August 16, 2026. Market data as of the August 14, 2026 close.



PUBLIC ARTICLE. Sarah, David, the chief financial officer, and Meridian Systems are illustrative composites. Factual claims about institutions, studies, and companies are sourced in the notes. The market analysis in Part V is an exploratory analysis of publicly available information, not a pre-registered valuation study or investment recommendation. Evidence cutoff August 16, 2026; market data principally as of August 14, 2026.

Who Can Afford Friction? How AI Reprices Economic Power

Sarah, David, the chief financial officer, and Meridian Systems are illustrative composites, assembled from common patterns rather than drawn from any real person or company. Market data is as of August 14, 2026 unless stated otherwise. Every factual claim about institutions, studies, and companies is sourced in the notes.

Part I. The Friction Economy

Two problems on the same morning

AI is making it extraordinarily cheap to produce things that look finished. An answer. An analysis. A contract. A working piece of software. A legal argument. A financial recommendation. A challenge to somebody else's conclusion.

But making something easier to produce does not make it easier to trust, and in some cases the opposite is true. As the cost of producing claims, arguments and counterarguments falls, more of them get generated, examined and contested. The scarce thing starts to move away from producing an answer and toward determining which answers deserve to be relied upon. By trust I do not mean whether AI feels trustworthy. I mean something more operational: **reliance**, being willing to let an answer support a consequential action.

I think one of the best ways to understand why that shift is happening is through something much more familiar: friction. Consider two people who discover problems on the same morning.

Sarah is buying a house. Two weeks before closing, her lender tells her something on her credit file has moved her rate. She pulls the report and finds an account she does not recognize, reported delinquent, from a lender she has never used.

She has a right to dispute it. That right is not the constraint. The constraint is everything standing between her and the exercise of it: understanding what the entry says, reconstructing four years of her own financial history, identifying which bureau and which furnisher, the company that supplied the information, to contact, learning what a dispute letter must contain to be taken seriously, assembling documents she may not have kept, and persisting through a process that will very likely tell her, on first pass, that the information has been verified as accurate.

The same morning, David is reading a competitor's product documentation. He founded a fourteen-person software company. His engineer has flagged that a much larger company appears to be doing something uncomfortably close to a method his company invented, built, and patented four years ago.

He also has a right, and his patent is real. His constraint is not the right either. It is the cost of finding out whether the suspicion is worth anything: mapping the competitor's implementation against his claims, gathering evidence he cannot obtain from outside, and buying enough specialist time to learn whether he has a case or an expensive misunderstanding. Answering *should I even look into this* has historically cost more than a company his size can casually spend. And answering it is not the same as knowing whether the answer will support what it would cost him to act on it.

Sarah wants to remove friction. David wants to exercise a right that creates friction for someone else. They are not opposites. They are the same economics viewed from two sides. And artificial intelligence lowers the cost of both.

That is the argument of this essay. **AI is not principally a friction-elimination technology. It is a technology for repricing friction.** It makes some barriers cheap to cross, others cheap to build and enforce, and leaves a residue it barely touches. The question is not whether friction falls, but who was paying for it, who benefits when the price changes, and what becomes scarce afterward.

Put more simply: **AI is not making the economy frictionless. It is changing who can afford friction.** And underneath that, the reason it matters: **getting an answer is becoming cheaper. Knowing when it is good enough to act on is not.**

I am not a detached observer of this. I build companies, I own equity, and I invest, and I have organized all three around the argument below. I return to it at the end, and everything in between should be read with that interest in view.

Three questions to carry through this essay

1. What friction is AI changing, and for whom?

2. Can the organization adapt before the economics move?

3. How much of that change is already in the price?

Friction was economic long before AI

Let me be precise about what is old here, because the old part is well established and I am not claiming it.

Ronald Coase argued in 1937 that firms exist at all because using the price mechanism is itself costly, and that a firm expands until the cost of organizing one more transaction internally equals the cost of buying it in the market.¹ George Stigler put information on the same footing in 1961: search has declining marginal return, so buyers rationally stop before finding the lowest price, and price dispersion survives as a measure of ignorance in the market.² Christopher Sims formalized scarce attention: an agent whose actions reach the world only through a finite-capacity channel responds in a delayed, smoothed, noisy way with no physical rigidity present at all.³ Paul Klemperer showed switching costs make products that were identical before purchase into differentiated ones after it, changing competition itself.⁴ Moynihan, Herd and Harvey decomposed the burden citizens face into learning, compliance, and psychological costs, and argued the level of each is frequently a political choice.⁵

¹Coase, R. H. (1937). "The Nature of the Firm." *Economica*, 4(16), 386–405. <https://doi.org/10.1111/j.1468-0335.1937.tb00002.x>. The phrase "transaction costs" is not Coase's here; he writes of the cost of using the price mechanism. The Coase theorem comes from "The Problem of Social Cost" (1960).

²Stigler, G. J. (1961). "The Economics of Information." *Journal of Political Economy*, 69(3), 213–225. <https://doi.org/10.1086/258464>

³Sims, C. A. (2003). "Implications of rational inattention." *Journal of Monetary Economics*, 50(3), 665–690. [https://doi.org/10.1016/S0304-3932\(03\)00029-1](https://doi.org/10.1016/S0304-3932(03)00029-1). Rational inattention is not bounded rationality in the heuristics sense; agents optimize fully subject to a processing constraint.

⁴Klemperer, P. (1987). "Markets with Consumer Switching Costs." *Quarterly Journal of Economics*, 102(2), 375–394. <https://doi.org/10.2307/1885068>. Net price and welfare effects are ambiguous: switching costs can raise later prices while intensifying early competition for share.

So friction is a real economic object with structure, it shapes market outcomes, and its distribution is often deliberate. None of that is mine. Here is what I think is new. Those literatures largely ask what friction *does* to economic behavior, holding the price of overcoming it roughly fixed. The AI question is what happens when that price falls sharply, unevenly, for some actors before others, in a period of a few years. That is not a refinement of transaction-cost economics. It is a shock to one of its inputs, and shocks to inputs redistribute rents.

Four kinds of friction

The reflex in technology is to treat friction as waste and its removal as progress. That produces bad strategy, because it collapses four different things into one word.

Type	What it is
Waste friction	Cost that produces little protective value for anyone, often surviving because no one had a reason to remove it
Asymmetric friction	Cost one side navigates far more cheaply than the other, whether or not anyone designed it that way
Protective friction	Deliberate cost imposed to achieve verification, authentication, safety, deliberation, or due process
Strategic friction	Barriers arising from legitimate ownership, exclusivity, rights, scarcity, or regulated permission

Cass Sunstein's work on sludge is useful precisely because it does not say what it is usually quoted as saying. He defines sludge, meaning excessive or unjustified procedural burden, as *excessive or unjustified* friction, then devotes a section to where burden is justified: countering impulsivity, protecting privacy and security, acquiring necessary information, establishing eligibility.⁶ The recommendation is an audit, not an amputation.

"All friction is exploitative" is false. Identity verification before a wire transfer is friction. A cooling-off period is friction. Discovery is friction. Remove all of it and you get an economy that is trivially cheap to defraud. The question is not whether friction exists, but what it does, who pays it, and what happens when its price changes.

A working definition: **business friction is the aggregate cost separating an actor from an economically, legally, contractually or technically available outcome.** Decomposed:

$$F = \text{search} + \text{knowledge} + \text{expertise} + \text{procedure} + \text{time} + \text{evidence} + \text{coordination} + \text{persistence} + \text{verification} + \text{authority}$$

Two things about that list. Read it as the drivers of one quantity, the fully loaded expected cost of reaching a named outcome, rather than ten independent primitives that happen to add up: expertise can be bought as

⁵Moynihan, D., Herd, P., & Harvey, H. (2015). "Administrative Burden: Learning, Psychological, and Compliance Costs in Citizen-State Interactions." *Journal of Public Administration Research and Theory*, 25(1), 43–69. <https://doi.org/10.1093/jopart/muu009>. Conceptual and case-based; does not provide a causal effect size for burden on take-up.

⁶Sunstein, C. R. (2022). "Sludge Audits." *Behavioural Public Policy*, 6(4), 654–673. <https://doi.org/10.1017/bpp.2019.32>. Published online January 6, 2020; print issue October 2022. The paper contains a dedicated section on when administrative burdens are justified.

time, persistence and coordination consume time, and each unit of cost is assigned once. The cost of the tooling itself stays inside the ledger. Model fees, checking the output and cleaning up its errors are per-attempt costs; integration, data preparation and privacy controls are one-time. All count against the friction AI removes, but not in the same units, so what matters is the net change in whichever of the two you are deciding about, never a blended figure. Gross task acceleration is not the quantity at all. Keep the fixed and the marginal apart, because they are not in the same units and cannot be added. Call them **F_fixed**, currency paid once to acquire the capability, and **F_marginal**, currency *per attempt*. Whether to adopt at all turns on F_fixed against the surplus you expect across every attempt you will make. Whether a particular \$18.73 dispute is worth pursuing turns only on F_marginal. To combine them you must amortize over a *stated* number of attempts, $F_{\text{marginal}} + F_{\text{fixed}} / N$: a \$240 subscription makes a \$2 dispute cost \$242 at $N = 1$ and \$22 at $N = 12$. The answer moves an order of magnitude on a number most analyses never write down. That said, the last two terms are the ones to watch. My thesis is that AI compresses the first eight costs much faster than it compresses **verification** and **authority**, and that this asymmetry generates most of what follows.

The economy of rational surrender

Start with the simplest case. A charge appears for \$18.73. You do not recognize it. You are fairly sure it is wrong.

Whether you challenge is not a question of justice. It is arithmetic. Call it the **Minimum Viable Challenge**:

$$p \times R > F_c$$

where R is what you recover, p is your probability of success, and F_c is the *marginal* cost of making this one attempt: the time, attention and unpleasantness it takes, plus any per-use cost of a tool. It is not the cost of having the capability at all, which is one-time, in different units, and governs a different decision. The same inequality recurs throughout this essay in different clothes: a probability times a value, set against the friction of getting it.

In plain English: a claim is worth pursuing when its expected benefit exceeds everything it costs to pursue it. Note what F_c is: the *marginal* cost of this one attempt, the hour, the attention, the two follow-ups, the per-use cost of any tool. Charging an annual subscription against a single \$18.73 dispute would be a category error. At a sixty percent chance of getting the charge reversed, expected recovery is \$11.24. If the dispute takes an hour of attention plus two follow-ups, call it an hour and a half, it clears only for someone valuing that bundle below \$11.24 in total, under \$7.50 an hour. At two hours the break-even falls to about \$5.60. Both are well below what most people's time is worth, which is why the dispute usually does not clear. What share of the population that is, I cannot tell you and will not pretend the model knows: it depends on how long the process takes, and the process is designed by the party you are disputing with. When it does not clear, declining the dispute is economically rational even though the underlying claim is valid. You are not being lazy. You are correct.

Now aggregate. Every rational decision not to challenge is also a decision that value stays where it landed. Multiply by millions of consumers and hundreds of millions of small discrepancies and you get **an economy of rational surrender**, in which a nontrivial amount of value is allocated not by contract or merit but by the fact that pursuing it costs more than it is worth.

This does not necessarily self-correct. Gabaix and Laibson showed shrouded attributes persist even in highly competitive markets with costless advertising, because a firm that unshrouds a rival's hidden fees mostly

teaches that rival's customers to avoid the fee rather than switch providers, so no firm profits from educating.⁷

I want to be careful about intent, because this is where the argument usually goes wrong. I am not claiming financial services or telecommunications deliberately engineered opacity to extract from confused customers. Some firms have; the FTC's staff report on dark patterns catalogues eight categories including obstruction, sneaking and interface interference, and the OECD reached similar conclusions internationally.^{8, 9} But the argument does not need intent and is stronger without it. **Asymmetric friction does not require a designer.** A complex product sold at scale to individuals produces cost asymmetry automatically, because the seller spreads process knowledge across millions of transactions and the buyer across one.

Which is why the AI move matters. AI does not need to increase the value of an \$18.73 claim. It only needs to reduce F_c , in exactly the terms that were doing the work: noticing, understanding, drafting, tracking, persisting.

I think the marginal cost of vigilance is collapsing. When vigilance was expensive, you spot-checked. When it is nearly free, the rational policy changes: audit everything you can afford to resolve. That last clause is not a hedge. Detection getting cheap does not make resolution cheap, and an audit that surfaces more exceptions than you can work is a queue, not a recovery.

AI may institutionalize the individual

Return to Sarah. What changes is not her rights. It is the cost of operating them.

A capable personal AI system could, in principle, maintain a record of her accounts, reconcile documents against each other, flag the contradiction between a furnisher's claim and her payment history, identify which statutory process applies, draft a dispute containing the elements that make it hard to dismiss, track deadlines, and prepare a package for a lawyer.

Note the verbs. None of them is *decide*. The AI does not adjudicate her claim. The narrower claim: individuals can now acquire something resembling institutional memory, procedural fluency and cheap on-demand expertise, capabilities previously concentrated in organizations because only organizations could amortize them.

AI may institutionalize the individual. Not to parity: a person with a good AI system is not a company with a compliance department and the ability to outlast you. But the gap narrows fastest where it was widest, in the low-value, high-nuisance disputes institutions have historically won by attrition.

Disputes at this level are not hypothetical: in January 2025 the CFPB ordered Equifax to pay a \$15 million penalty over its handling of credit-report disputes, alleging it disregarded documents consumers submitted, allowed deleted information to reappear, and sent letters telling consumers in the same communication that

⁷Gabaix, X., & Laibson, D. (2006). "Shrouded Attributes, Consumer Myopia, and Information Suppression in Competitive Markets." *Quarterly Journal of Economics*, 121(2), 505–540. <https://doi.org/10.1162/qjec.2006.121.2.505>. Conditional on myopic consumers and cheaply avoidable add-ons; not a claim that markets never self-correct.

⁸Federal Trade Commission, *Bringing Dark Patterns to Light* (September 2022). A staff report, not a Commission rule. <https://www.ftc.gov/reports/bringing-dark-patterns-light>

⁹OECD (2022), *Dark Commercial Patterns*, OECD Digital Economy Papers No. 336. <https://doi.org/10.1787/44f5e846-en>

an item had been verified as accurate and that it had been deleted.¹⁰ And the volume is real: the CFPB received 6,635,400 complaints in 2025 and sent about 90 percent of them to companies for response.¹¹

Now the caution. **Complaint volume is not validated wrongdoing.** Six million complaints are six million assertions, and the agency that collects them said so. In June 2026 the Bureau described its own complaint system as distorted, noted that credit and consumer-reporting complaints had grown from more than 150,000 in 2019 to more than five million in 2025, attributed that distortion to credit repair organizations misusing the process, to social media promotion, and to new technologies including AI tools, and concluded it could not rely on the portal data as a reliable reflection of market conditions or consumer experiences.¹²

It is not evidence that AI caused the increase; the Bureau names suspected contributors, not a causal estimate, and neither it nor I have established one. But it is a real institution reporting that cheap high-volume submission created a verification problem for the institution receiving it.

The interpretive point stands regardless:

The law did not change. The cost of making the law operational for one person did.

Key takeaway. AI changes the economics of exercising rights and navigating processes, not merely the cost of producing work.

Ask yourself. What outcome do your customers, employees or counterparties abandon today because pursuing it is not worth the effort? Who navigates that same process more cheaply?

Part II. The Other Side of Friction

The mirror image: activating legitimate rights

Everything above reads as consumers gaining ground against companies. That is the less interesting half. Consumer finance shows AI making friction cheaper to *overcome*; intellectual property shows it making legitimate friction cheaper to *exercise*.

Return to David. A right to exclude is not self-executing, and the gap between holding it and using it is measured in money. The Government Accountability Office, reporting in December 2024 on third-party funding of patent litigation, cited an AIPLA survey putting the median cost of litigating a patent case through trial at more than \$3 million depending on damages sought.¹³ For a fourteen-person company, that does not describe a difficult decision. It describes an unavailable one.

¹⁰CFPB, "CFPB Orders Equifax to Pay \$15 Million for Improper Investigations of Credit Reporting Errors," January 17, 2025. These are the Bureau's allegations and the terms of a consent order.

¹¹CFPB, *2025 Consumer Response Annual Report*, published March 31, 2026. 6,635,400 complaints received; approximately 5,984,100 (90 percent) sent to companies.

¹²CFPB, "The CFPB is Correcting Flaws to Restore Integrity and Utility to the Consumer Complaint System," June 24, 2026. The Bureau names suspected contributors, not a causal estimate, and neither it nor I have established causation.

¹³U.S. Government Accountability Office, *Intellectual Property: Information on Third-Party Funding of Patent Litigation*, GAO-25-107214, published December 5, 2024. The median cost figure is attributed by GAO to a 2023 AIPLA survey. Both the access argument and the weak-claims concern are stakeholder views collected by GAO, not GAO findings of fact.

State the parallel formally. The **Minimum Viable Enforcement**:

$$E[B_E - C_E] > 0$$

where B_E is the benefit realized under each possible outcome, before enforcement cost, and C_E is the corresponding cost of investigating and enforcing under that same outcome. The expectation weights those outcomes by their probabilities; neither term is itself an average, and neither is the payoff you would receive if you won. The outcomes to weight are recovery, licensing, settlement, deterrence and strategic effect, and where they are alternatives rather than complements their probabilities must be jointly feasible and cannot sum above one; adding them independently inflates the left side by just enough to make marginal cases look worth filing. Keeping benefit and cost as separate distributions matters: calling the left side "net" and then comparing it to cost would subtract the cost twice. I write it this way rather than as one probability times a sum of outcomes because those outcomes are often alternatives rather than complements: winning damages can foreclose the licensing path, and deterrence accrues without formal success. There is a second decision hiding inside the first, and it is the one David actually faces. Investigating is cheap relative to enforcing, and what it buys is the right to decide later. Paying to find out can be rational even when filing is not, because the answer tells you whether to proceed or drop it. Note that the inequality above does not capture that sequence; it is a one-shot threshold, and the option to learn first is worth something it does not price.

One more thing the inequality does not price, and it is the one that matters most for a fourteen-person company. A positive expected value is not an affordable bet. Suppose the payoff if you prevail is \$10 million, the chance of prevailing is forty percent, and the cost is \$3 million whichever way it goes. Then B_E , the expected benefit, is \$4 million, the expectation $E[B_E - C_E]$ is positive by a million dollars, so the expected-value condition is satisfied, while the actual distribution is a sixty percent chance of spending three million dollars and recovering nothing. For David that is not a disappointing outcome, it is the end of the company. The threshold is a necessary condition, not a sufficient one, and the prior question is whether you can survive the branch where you lose. This is exactly why cheaper enforcement changes who can act and not merely how much acting costs: the constraint that binds small holders is not the expectation, it is the variance they cannot absorb. In plain English: a valid right stays dormant when discovering, proving and exercising it costs more than it is expected to be worth, or when the losing branch would be fatal. It exists on paper and does nothing.

How much sits below the line? The GAO report is instructive because it is not about AI. It describes what happens when a *different* constraint on F_E relaxes: university officials and inventors told GAO that third-party funding lets resource-constrained owners file suits they otherwise could not. ¹⁴ **Relax the cost constraint around enforcement and dormant rights become active.**

The same report cuts the other way. Technology companies told GAO many funded cases involve weak infringement claims, and that defendants incur real costs even on patents they expect to be invalidated. ¹⁵ Those are stakeholder views, not GAO findings, but they apply with full force. **Cheaper enforcement is not automatically good.** Lowering F_E activates meritorious and meritless dormant claims at once. Anyone who tells you AI will democratize enforcement without also cheapening harassment is selling something.

Now the AI part, carefully. The best current evidence is a 2026 randomized trial by Schwarcz and colleagues: upper-level law students, 137 of whom completed at least one task, randomly assigned to a retrieval-

¹⁴U.S. Government Accountability Office, as cited above.

¹⁵U.S. Government Accountability Office, as cited above.

augmented legal tool, a reasoning model, or no AI across six tasks. Significant gains appeared on five of six, ranging from 50 to 130 percent, with improved quality.¹⁶ What it does not say matters more. It does not say practicing lawyers gained 50 to 130 percent; the participants were students on stylized exercises. It does not say hallucination is solved; the reasoning model produced eleven fabricated or mis-stated citations against three for the retrieval tool and four for the unassisted control group. And the contract-drafting task showed no meaningful improvement.

AI does not make legal expertise free. It changes the shape of the workflow, compressing research, drafting, organizing and triage while leaving expensive judgment where it was, and arguably moving it later.

David does not need a verdict from a machine. He needs to turn an unanswerable question into an answerable one. The boundary: **AI cannot determine infringement.** What it can reduce is the cost of searching, mapping and triaging the questions that may warrant qualified review, and so the cost of finding out whether the specialist call is worth making.

The institutions are moving the same way. The Patent and Trademark Office's 2024 guidance states directly that AI tools have the potential to lower the barriers and costs of practicing before the Office, while insisting professional obligations are undiminished and that simply relying on the accuracy of an AI tool is not a reasonable inquiry.¹⁷ The Office also ran an AI-assisted pre-examination prior-art search pilot from October 2025 until it closed to new petitions after June 1, 2026, with modest uptake.¹⁸ That is evidence computational search is entering the examination side, not evidence it produces better patents.

AI changes IP strategy without abolishing patent law

If exploration gets cheaper, patent strategy changes before patent law does.

Drafting a strong specification has always been limited by how much ground a team can afford to cover: embodiments, alternative architectures, prior art, the ways a competitor might implement the same idea differently. Those are search and analysis costs, exactly what AI compresses. I expect well-run companies to explore that space far more exhaustively, though that is a hypothesis, not a result.

The goal is not a patent that is impossible to design around; that is wrong on the law and on the economics. The relevant question is a distance question: **how far must a competitor move from its preferred implementation to avoid the protected territory, and what does the detour cost?** A patent forcing a rival into an inferior architecture has created real economic separation even though a design-around exists. Call it design-around distance. It is continuous, not binary.

Two constraints keep this honest. Written description and enablement are separate statutory requirements under 35 U.S.C. 112(a); an invention may be described without the disclosure being enabling, and a disclosure may be enabling without describing the invention.¹⁹ You cannot claim territory you have not adequately

¹⁶Schwarcz, D. B., Manning, S., Prescott, J. J., Barry, P., Cleveland, D. R., & Rich, B. (2026). "AI-Powered Lawyering." *Journal of Law & Empirical Analysis*, 3(1), 220-250. <https://doi.org/10.1177/2755323X261427048>. RCT run October 2024; 153 enrolled, 137 completed at least one task, 125 completed all six. The 50 to 130 percent figure is the range of significant gains across tasks and tools, not a uniform speedup; per-task time reductions in the text are roughly 20-28 percent. Results do not generalize to practicing lawyers on live matters.

¹⁷USPTO, "Guidance on Use of Artificial Intelligence-Based Tools in Practice Before the USPTO," 89 Fed. Reg., April 11, 2024.

¹⁸USPTO, Artificial Intelligence Search Automated Pilot Program (ASAP!). Launched October 20, 2025; expanded April 16, 2026; closed to petitions after June 1, 2026. As of April 2026, 169 petitions filed and 76 granted. Institutional adoption is not evidence of improved patent quality.

¹⁹MPEP 2161, "Three Separate Requirements for Specification Under 35 U.S.C. 112(a)."

taught. And the competitor has the same tools: if AI makes it cheaper to map a claim's boundaries from inside, it makes it cheaper from outside too.

When everyone can build

The same repricing is visible in software production, where the friction story gets its largest-sample measurement.

A 2026 NBER working paper by Demirer, Musolf and Yang combines data on more than 100,000 GitHub developers with usage telemetry in a matched event study. Autocomplete, interactive coding agents, and autonomous coding agents each significantly increase commits, with cumulative effects of 40, 140 and 180 percent. Then the finding that matters: the gains attenuate sharply across the production hierarchy. The 180 percent effect on commits falls to 50 percent at the level of projects and to 30 percent for actual releases.²⁰ A cross-check across app marketplaces found a moderate increase in new apps and no increase in total usage.

Read that as a friction measurement rather than a productivity headline. The study estimated roughly a 180 percent increase in commits and roughly a 30 percent increase in releases. Those are two different activity measures, and neither one measures a pass-through from code to economic value. My narrower inference is that downstream constraints absorb much of the increase in coding activity, so the binding constraint was never only typing.

That attenuation is one observational study, not a law, but it illustrates a pattern I think is general. **AI does not remove bottlenecks. It relocates them.** A process is only as fast as its binding constraint, and accelerating a non-binding step reveals which step was actually binding.

Two studies keep the picture honest. Brynjolfsson, Li and Raymond, studying 5,172 customer support agents, found a 15 percent average increase in issues resolved per hour, concentrated among less experienced workers, while the most experienced saw small gains in speed and small declines in quality.²¹ Dell'Acqua and colleagues, in a field experiment with 758 management consultants, found that inside the AI's capability frontier participants completed 12.2 percent more tasks, 25.1 percent faster, at higher quality, while on a task deliberately chosen to sit outside it, AI users were substantially *less* likely to reach the correct answer: 60.0 and 70.6 percent against a control group's 84.5 percent.²²

That is not a technology that uniformly reduces the cost of thinking. It reduces some costs sharply, leaves others alone, and raises a few.

The effect on defensibility is real regardless. **Implementation difficulty alone is becoming a weaker source of defensibility**, and any moat that amounted to *we built it and replicating it would annoy you* is worth less. When everyone can build, what matters is what everyone cannot simply copy: distribution, trust, proprietary data, regulated permission, exclusive relationships, and rights.

²⁰Demirer, M., Musolf, L., & Yang, L. (2026). "Writing Code vs. Shipping Code." NBER Working Paper No. 35275, May 2026. Matched event study on observational telemetry from more than 100,000 GitHub developers. Unpublished working paper, not peer reviewed. Reports activity and output counts, not code quality or economic value.

²¹Brynjolfsson, E., Li, D., & Raymond, L. R. (2025). "Generative AI at Work." *Quarterly Journal of Economics*, 140(2), 889–942. <https://doi.org/10.1093/qje/qjae044>. Staggered-rollout study at a single software firm, not a randomized experiment.

²²Dell'Acqua, F., McFowland III, E., Mollick, E., Lifshitz, H., Kellogg, K. C., Rajendran, S., Kraymer, L., Candelon, F., & Lakhani, K. R. (2026). "Navigating the Jagged Technological Frontier." *Organization Science*, 37(2), 403–423. <https://doi.org/10.1287/orsc.2025.21838>. The 758 participants were Boston Consulting Group consultants on purpose-built exercises rather than live client work, using GPT-4 as of 2023. The outside-frontier result rests on one task deliberately engineered to sit beyond the model's capability.

Traditional moats are not being destroyed. They are being repriced, and some upward.

Key takeaway. The same technology that makes friction easier to overcome makes rights, legitimate and otherwise, cheaper to exercise.

Ask yourself. What rights does your company already own but rarely investigate because the cost is too high? Which rights held against you become active when the other side's cost falls?

Part III. How I Analyze Opportunities

The Friction Balance Sheet

This is an inventory, not an accounting balance sheet. The four quadrants hold different kinds of thing measured in different units, entries can sit in two quadrants at once, and there is no net position to compute. Every company holds friction on both sides of a ledger. Most have never written it down, which means they do not know their own exposure.

Friction assets	Friction liabilities
Legitimate IP, exclusive contracts, trusted brand, distribution, regulated permissions, proprietary data, physical scarcity: positions whose value <i>rises</i> when the cost of exercising them falls	Margins that depend on a counterparty finding it too costly to compare, inspect, switch, remember, or object
Friction optionality	Friction exposure
Rights you already own but have never been able to monitor or enforce economically	Rights <i>others</i> hold against you that were dormant only because exercising them was expensive

The bottom two quadrants are the ones nobody tracks, and the ones AI moves first.

Consider a chief financial officer at a mid-size manufacturer with roughly 1,800 supplier agreements containing price adjustment mechanisms, rebate thresholds, service level credits, most-favored-customer clauses and audit rights. Every one is a real, negotiated, paid-for right. Almost none are monitored, because checking 1,800 contracts against actual invoices would cost more than the recoveries. They are dormant for the same reason Sarah's dispute and David's patent are dormant: the threshold is not met, one clause at a time.

AI does not create those rights. It changes the cost of converting them into realized value. **Latent rights** are entitlements you hold that are real but not currently economical to exercise. **Latent obligations** are the mirror: entitlements held *against* you, which will not stay dormant if the other side's costs fall too.

That symmetry is the point. **"AI empowers consumers, therefore companies lose" is wrong, and so is its opposite.** Both sides are being armed. The question is which entries on your own balance sheet move, in which direction, and how fast.

Five tools I use

I use five tools when I look at a business through this lens. They are not established metrics or accounting standards, and I would resist anyone treating them that way. Their value lies in refusing to let five different questions collapse into one score.

Friction asymmetry ratio. Who bears it? For a named interaction o , $FAR(A:B,o) = F(A,o) / F(B,o)$: the fully loaded friction borne by one party over that borne by the other, for the same outcome, on the same cost basis, over the same period. In plain English, how much harder is it for A than for B to accomplish the same contested thing? At 1 the two sides bear the same cost, which is not the same as bearing the same burden: FAR is a ratio of absolute money, and \$300 is not the same imposition on a household as on a ten-billion-dollar institution. If burden is the question, divide each side by a stated measure of capacity to bear it and say which measure you used. Well above 1, one side carries the relationship. It is meaningless in the aggregate and must always be stated for a specific interaction. "Banking has an asymmetry of 8" is not a statement about anything. "For challenging this category of account error, the consumer's fully loaded burden per dispute may substantially exceed the institution's fully loaded burden per case" is, provided you say over how many attempts each side's fixed cost is spread, because the institution spreads its process cost across millions of cases and the consumer across one. Those are two different values of N in the amortization above, and a FAR computed without stating both is measuring the undeclared N as much as the asymmetry. Both legs must be on the same basis; comparing one side's fully loaded cost to the other's marginal cost measures the accounting convention, not the asymmetry.

Friction dependency ratio. How much value depends on friction staying expensive? Divide value dependent on compressible friction by the value under analysis, whether revenue, gross profit, retention, a product line, or enterprise value. Numerator and denominator must be the same measure: gross profit against gross profit, not revenue exposed over enterprise value. It behaves as a share only when the exposed value is genuinely a subset of a positive base; with a loss-making segment inside the total it can exceed one, and for a loss-making company it is not defined at all. Two companies each earning \$100 million of gross profit can look identical while one earns it from scarcity, technology and distribution and the other from inertia, obscure pricing and switching cost. The diligence question I actually ask: if every customer had a perfect analyst and negotiator tomorrow morning, how much of this company's economics would still be there on Friday?

Friction conversion. How much of what we already own are we collecting? Divide value realized from identified exercisable rights by identified economically available value. The denominator must be bounded, frozen for the period, and on the same gross or net basis as the numerator, or it becomes a fantasy number and an analyst improves conversion simply by changing what is admitted to it. Measure one cohort forward: value identified at the start of the period against value realized from that same set by the end. Scoring what you captured last year against what a review found this year manufactures a failure rate out of hindsight. Return to the CFO: a review identifies \$400,000 in service credits, \$250,000 in unused price protections, \$150,000 in rebates and \$200,000 in recoverable overcharges. That is a \$1 million cohort, frozen at the start of the period. If monitoring realizes \$700,000 of that same cohort by period end, conversion is 70 percent, and not one new right was acquired. The comparison the rule forbids is against an opportunity set nobody had enumerated at the time. Then run it in the mirror, or the tool degrades into a collections exercise: what proportion of *our* obligations can counterparties now detect and enforce?

Friction beta. How sensitive is value to a change in friction? This is not CAPM beta; I borrow the word only for sensitivity. Define it from the elasticity of an economic outcome with respect to a named friction: β_F is *minus* that elasticity, which is to say minus the percentage change in value over the percentage change in friction. The minus sign is there so the reading is intuitive, and it does mean β_F is not itself the elasticity but its negative. It is also local: it describes a small move in F , and most of the frictions in this essay are moving a long way. Positive beta, value rises as friction falls, which is the patent portfolio when monitoring gets cheap. Negative beta, value falls as friction falls, which is the retention line when switching gets easy. One beta per named friction, everything else held still, and betas taken with respect to different frictions do not add: to combine them you would need a forecast of how far each friction actually moves. A business that gains when contracting, compliance or enforcement gets cheaper has positive friction beta. A business whose margin rests on inattention, difficult switching, or opacity has negative friction beta. This is a better question than "is this company exposed to AI," because it demands an answer to: exposed to *which* friction, in which direction, over what horizon.

Friction frontier. What becomes worth doing? Action is rational when $p \times V$ exceeds F . At the boundary $p \times V = F$, so the smallest worthwhile opportunity is $V^* = F / p$. Lower F and V^* falls with it, holding p constant. Be careful with p : the p in this formula is the probability of *succeeding*, not of *spotting* a candidate. Better screening raises the number of candidates you examine rather than lowering the threshold. It also changes which claims enter the pool, and so changes the average success probability among them; which way that runs is an empirical question, not something the formula settles. *Do not spend five hours contesting a \$40 error becomes the system flags it, documents it, and files it. Nobody audits 1,800 contracts becomes somebody does, nightly.* The important effect is not that existing work gets cheaper. It is that whole classes of work start existing.

Two rules keep this from becoming storytelling. Every measurement names whose cost, against which counterparty, toward which outcome, in which direction, over what period. "Consumer-to-lender dispute friction, procedural, falling, over three years" is meaningful; "financial services friction" is not, and a number attached to it can be chosen after the conclusion. And no false precision: most applications will be qualitative or estimated from stated assumptions, a few measured. Say which. A friction beta of 1.73 asserted on inputs that cannot support two decimals discredits everything around it.

One worked case, because five definitions prove nothing. Meridian Systems is a hypothetical subscription software company with difficult cancellation, real proprietary technology, long customer contracts, a vendor portfolio, and ordinary support costs. *Asymmetry*: cancelling costs the customer far more than processing the cancellation costs Meridian. *Dependency*: some share of retention exists because leaving is annoying rather than because staying is better, and that share is the exposed share. *Conversion*: Meridian is almost certainly realizing a fraction of its supplier rights and watching almost none of its intellectual property. *Beta*: retention has negative friction beta, because easier switching costs it customers, while the patent portfolio has positive friction beta, because cheaper monitoring makes the same rights worth more. *Frontier*: auditing every vendor agreement, watching the market for infringement, and reviewing every disputed transaction each cross from uneconomic to economic, but not together. Contract and transaction review turn on search and reading cost, where compression is fastest; infringement watching still ends in a determination of infringement, which is verification and authority, the slowest terms in the ledger. Expect the reading-heavy thresholds to fall first and by more. And a frontier that moves creates nothing unless opportunities sit in the space it opens.

Notice what happened. The same company is short friction in one place and long friction in another, simultaneously. Borrowing the trading sense: you are long a friction when your economics improve if it persists,

and short it when they improve if it falls. **There is no single scalar called a company's friction exposure.** There is a portfolio of friction exposures, each with its own direction, size and speed, and no admissible way to net them: an annual margin flow and the present value of a contingent right are not the same kind of quantity, and subtracting one from the other produces a number that means nothing. Look at the entries, not at a total. A company can be quiet on every summary measure you have while individual entries move violently in opposite directions, which is precisely why the summary measure is the wrong thing to watch.

Key takeaway. A company can be long friction in one part of its economics and short friction in another, at the same time.

Ask yourself. If every customer had cheap expertise and perfect memory tomorrow, which part of your margin would still be there? Which latent rights or obligations would suddenly matter?

Part IV. Can the Company Change?

The second lens: can the organization change?

The friction lens tells you where economic value is likely to move. It does not tell you who captures it. A company can have a vulnerable business model and survive by recognizing the change early, cannibalizing its own economics, and reallocating capital. A company can have excellent assets and lose because it cannot change fast enough to use them.

Structural exposure is not adaptive capacity, and knowing change is necessary is not being able to execute it.

The analogy is not that the automobile replaced the horse. It is that motorization repriced an entire interlocking system. American farm horses and mules peaked at 26.7 million head in 1918; cropland used to feed draft animals fell from roughly 93 million acres in 1915 to about 5 million by 1960.²³ Motor vehicle registrations went from 468,500 in 1910 to 26.7 million in 1930, which by coincidence is almost exactly the horse peak twelve years earlier.²⁴ What moved was not one product category. It was the economics of distance, and with it roads, fuel, retail, real estate, logistics, insurance, consumer finance, agriculture and urban form.

That is the shape of a general purpose technology. Bresnahan and Trajtenberg define a general purpose technology, one used widely, improving over time, and enabling complementary innovation, by three features: pervasiveness, inherent potential for technical improvement, and innovational complementarities, and their central point is a coordination failure: arm's-length transactions between the technology and its users can produce "too little, too late" innovation.²⁵ Paul David's dynamo study makes the point historically:

²³Olmstead, A. L., & Rhode, P. W. (2001). "Reshaping the Landscape: The Impact and Diffusion of the Tractor in American Agriculture, 1910–1960." *Journal of Economic History*, 61(3), 663–698. <https://doi.org/10.1017/S0022050701030042>

²⁴Federal Highway Administration, *Highway Statistics Summary to 1995*, Table MV-200. The series is not monotonic; registrations fell between 1930 and 1935.

²⁵Bresnahan, T. F., & Trajtenberg, M. (1995). "General purpose technologies: 'Engines of growth'?" *Journal of Econometrics*, 65(1), 83–108. [https://doi.org/10.1016/0304-4076\(94\)01598-T](https://doi.org/10.1016/0304-4076(94)01598-T). The paper says nothing about asset prices or valuation, and does not identify AI as a GPT.

electrification's productivity payoff lagged its invention by roughly four decades because the gains required complementary reorganization.²⁶

The strategic question for a carriage company in 1910 was never whether the automobile was better than a carriage. It was whether the organization could become something other than the thing that had made it successful. And here the folk version is wrong. Carroll, Bigelow, Seidel and Tsai studied 2,197 American automobile producers from 1885 to 1981 and found the greatest longevity advantages among diversifying entrants accrued to bicycle and carriage manufacturers, with significantly lower death rates than other lateral entrants including engine manufacturers, attributed to assembly and coordination competence plus retail distribution experience.²⁷ Exposure for the carriage trade was near-total; conditional on attempting the transition, carriage firms were among the better adapters. Exposure and outcome are different variables.

AI is not simply human worker to AI worker. It is a change in the economics of cognition, and it therefore propagates into businesses far outside companies that sell AI, exactly as motorization reached businesses that never built a car.

Three failure modes, and the Change Gap

Organizations fail to adapt in three distinct ways, and conflating them produces bad analysis.

Recognition failure. The organization misunderstands what is changing. Tripsas and Gavetti's study of Polaroid is the canonical case, and it is routinely misremembered. Polaroid did not miss digital imaging; it built world-class digital capability early. What it could not escape was an inherited belief structure about where profit lived, which meant capability accumulated and commercialization did not.²⁸ The modern version: a company concludes "AI will reduce our customer service cost" when the real change is that its pricing architecture is about to become legible to every customer.

Change aversion. The organization could change but will not, because change threatens current economics. This is the Christensen and Bower mechanism: the internal resource-allocation process is captured by existing customers' revealed demands.²⁹ The constraint is willingness.

Change incapacity. The organization wants to change and cannot execute. Legacy systems, silos, slow decisions, talent gaps, technical and organizational debt. The constraint is capability.

Clark Gilbert's study of eight newspaper organizations separates the last two. He unbundles inertia into *resource rigidity*, failure to change investment patterns, and *routine rigidity*, failure to change the processes that use those resources. His finding is a sign reversal: a strong perception of threat **overcomes** resource rigidity while simultaneously **amplifying** routine rigidity, through contraction of authority, reduced experimentation,

²⁶David, P. A. (1990). "The Dynamo and the Computer: An Historical Perspective on the Modern Productivity Paradox." *American Economic Review (Papers and Proceedings)*, 80(2), 355–361. A historical analogy about one technology, not a predicted lag for AI.

²⁷Carroll, G. R., Bigelow, L. S., Seidel, M.-D. L., & Tsai, L. B. (1996). "The fates of de novo and de alio producers in the American automobile industry 1885–1981." *Strategic Management Journal*, 17(S1), 117–137. <https://doi.org/10.1002/smj.4250171009>. The authors also offer an exit-barrier interpretation: the carriage trade had already waned by 1910, so diversifiers had nowhere to retreat.

²⁸Tripsas, M., & Gavetti, G. (2000). "Capabilities, Cognition, and Inertia: Evidence from Digital Imaging." *Strategic Management Journal*, 21(10–11), 1147–1161. The paper refutes the "Polaroid missed digital" reading.

²⁹Christensen, C. M., & Bower, J. L. (1996). "Customer Power, Strategic Investment, and the Failure of Leading Firms." *Strategic Management Journal*, 17(3), 197–218. Incumbents typically developed the disruptive technology first; the failure is in commercialization funding.

and a focus on existing resources. Threatened sites raised online spending by 200 to 400 percent annually, and seven of eight still simply replicated the print newspaper online.³⁰ Money moved. Behavior did not.

Spending is resource mobility. It is not routine change.

Hannan and Freeman explain why this is normal rather than pathological: selection favors reliable, accountable organizations, reliability requires reproducing structure, and so selection favors inertia. Their definition is relative: structures have high inertia when reorganization is much slower than the rate at which conditions change.³¹ Teece, Pisano and Shuen stress that the countervailing capability is path-dependent, not available on demand.³² Tushman and O'Reilly, and Gilbert, converge on the same answer: structurally separated units.³³

I assess change capacity on eight dimensions, high, medium or low, never a composite number, and with demonstrated history weighted far above stated intention.

Dimension	The test
Recognition	Does management understand what is actually changing?
Decision speed	How long from evidence to committed action?
Cannibalization tolerance	Will it damage today's economics to hold tomorrow's position?
Capital mobility	Can money and people move off legacy activities?
Technical adaptability	Can the systems and data support a different operating model?
Organizational adaptability	Can roles, incentives and workflows actually change?
Execution capability	Do announced transformations happen?
Learning velocity	Can it recognize a wrong move and change again?

Return to Meridian. It scores high on recognition, because management can describe the seat-compression threat accurately. High on capital mobility, because it has cash and a small enough engineering organization to redirect. Low on cannibalization tolerance, because the pricing model that is most exposed is also the one the board's forecast rests on. Medium on execution, because it has shipped hard things before but never one that reduced its own revenue on purpose. That profile does not resolve into a score. It resolves into a prediction: Meridian will fund the response generously and implement it in a form that protects the existing price. An AI strategy slide is not change capacity. A capex commitment is not execution.

³⁰Gilbert, C. G. (2005). "Unbundling the Structure of Inertia: Resource versus Routine Rigidity." *Academy of Management Journal*, 48(5), 741–763. <https://doi.org/10.5465/amj.2005.18803920>. Inductive multiple-case study yielding propositions, not hypothesis tests; the 200 to 400 percent figures describe the sample rather than estimated effects.

³¹Hannan, M. T., & Freeman, J. (1984). "Structural Inertia and Organizational Change." *American Sociological Review*, 49(2), 149–164. A population-ecology argument about rates across populations, not a prediction about any individual firm.

³²Teece, D. J., Pisano, G., & Shuen, A. (1997). "Dynamic Capabilities and Strategic Management." *Strategic Management Journal*, 18(7), 509–533. A conceptual paper with no dataset.

³³Tushman, M. L., & O'Reilly, C. A. III (1996). "Ambidextrous Organizations." *California Management Review*, 38(4), 8–29. A conceptual managerial essay, not an empirical test that ambidextrous firms outperform. See also Leonard-Barton, D. (1992), "Core Capabilities and Core Rigidities," *Strategic Management Journal*, 13(S1), 111–125.

Hannan and Freeman's relative definition gives the closing move. The question is not whether a company can change but whether it can change faster than the economics are changing. A firm that needs three years to complete a transformation, in a market that demands the same scope of change every eighteen months, is moving at about half the required rate. The comparison only means anything for the same scope of transformation; two different-sized changes give a ratio of nothing. It may have excellent execution by historical standards and still lose. **Absolute change capacity is not adequate change capacity.** Adequacy is a ratio: organizational change velocity over the change velocity the environment demands, for the same scope of transformation. Above one the gap is closing, at one the gap is constant, below one it is widening. Note that this governs the direction of the gap, not its size. A firm three transformations behind and running at 1.4 is closing, and still loses if the horizon arrives first. Adequacy is a statement about the accumulated gap, which needs a starting position as well as a rate. I do not put a number on it unless both rates can be defined on the same basis, which is rare.

The gap between change required and change available is where opportunity concentrates. Incumbency cuts both ways: customers, data, capital, brand and regulatory knowledge are adaptation resources, while legacy revenue, incentives, politics and fear of cannibalization are adaptation resistance. Which dominates is the question, and assuming startups always win technological transitions is lazy.

Key takeaway. Exposure is not destiny. What matters is whether the organization can recognize, choose and execute change faster than its economics change.

Ask yourself. What has this organization actually cannibalized, reallocated or rebuilt before? Is management describing change, funding change, or demonstrating change?

One line matters more than the rest of this section. **One of the most dangerous mistakes in an AI transition is to mistake friction customers tolerate for value customers actually buy.** High retention, high margins and long relationships do not tell you which portion is value and which is switching cost. Historical financial performance cannot separate them. AI may make that separation tractable.

Part V. Price Is the Third Lens

Identifying a structural change is not identifying an opportunity. A company can be extraordinarily exposed to a friction about to collapse and still be a poor investment, because everyone else sees it too and the price already carries it. Economic insight is not investment edge. The edge exists only where the economics implied by the framework differ from those implied by the price.

Markets are not waiting for a theory of friction before reacting to AI. A Federal Reserve staff note describes financial markets as highly responsive to changes in the AI narrative, and a single week in February 2026 took roughly a trillion dollars off US software stocks.^{34, 35} Neither of those establishes sustained differentiation

³⁴Soto, P. E., Thieu, M., & Allen, J. S., "The AI Buildout and the Economy," FEDS Notes, Board of Governors of the Federal Reserve System, July 17, 2026. FEDS Notes represent the authors' views, not the Board's. The note also credits AI-related investment as among the largest positive contributors to quarterly GDP growth.

³⁵Singh, M., & Ahmed, S. I., "US software stocks slammed on mounting fears over AI disruption, lose \$1 trillion in week," Reuters, February 5, 2026. A discrete selloff event, not evidence of sustained multiple compression.

between businesses, which is the thing that would actually matter here; the second is one broad selloff rather than a dispersion. So the interesting question is not whether AI is priced. It is what, exactly, is priced.

Which forces a distinction most conversation about AI and valuation runs together. Three different things can be true of a company.

Awareness. Investors understand that AI materially affects this business.

Direct-effect pricing. The price reflects specific consequences: labor substitution, seat compression, pricing pressure, compute demand, cheaper information.

Full friction economics. The price reflects which friction is actually moving, how dependent the business is on it, which direction its exposure runs, whether the organization can adapt in time, what crosses the frontier next, which latent rights and obligations activate, and the long-run economics that follow.

Awareness and direct-effect pricing are visible in ordinary market commentary. I have no evidence the third is systematically priced, and I will not assume it is systematically unpriced either. That is narrower than the position I would prefer, and it is the one the evidence supports. It is also where reliance re-enters: acting on the difference requires knowing your model of it is good enough to bet on.

The reason this distinction is worth the trouble is that superficially similar AI-related price movements can rest on completely different economics. Consider credit scoring. In July 2025 the Federal Housing Finance Agency opened mortgage credit scoring to VantageScore 4.0 and then clarified that lenders delivering loans to Fannie Mae and Freddie Mac may use a single nationwide model, either Classic FICO or VantageScore 4.0, reversing a plan that would have required both; in April 2026 FHFA and HUD announced immediate acceptance of VantageScore 4.0 across Fannie Mae, Freddie Mac and FHA.³⁶ A position that had been effectively unavoidable in mortgage underwriting became one choice among two. Fair Isaac traded at \$1,085.78 on August 14, 2026, against a fifty-two-week high of \$1,998.01, roughly 46 percent lower.³⁷ Those price levels are public market data; reading the move as a consequence of the rule change rather than of something else is my own attribution, and no source here establishes it.³⁸

That case matters precisely because the mechanism is regulatory rather than technological. Friction gets repriced by regulation, competition, changing institutions and technology. AI matters because it can move several frictions unusually quickly and at once, not because every friction repricing is an AI event. A framework that treats every repricing as an AI event will misread the ones that are not, and will misprice the adaptation path for the ones that are.

³⁶Federal Housing Finance Agency, credit-score policy. On July 8, 2025 FHFA announced that VantageScore 4.0 could be used for loans delivered to the Enterprises; a July 15, 2025 clarification established that a lender may select a SINGLE nationwide model, Classic FICO or VantageScore 4.0, reversing the October 2022 plan that would have required delivery of both. On April 22, 2026 FHFA and HUD jointly announced that Fannie Mae and Freddie Mac accept VantageScore 4.0 for mortgage loans from all approved lenders effective immediately, completing implementation of the 2018 Credit Score Competition Act. <https://www.fhfa.gov/policy/credit-scores>. These sources establish the regulatory change and its timing. They do not establish the price levels, which are separately sourced, and they establish nothing about cause. They do not establish the magnitude or the cause of any movement in Fair Isaac's share price; attributing the repricing to this change rather than to something else is my own reading and is offered as such.

³⁷Fair Isaac Corporation (FICO) closing price \$1,085.78 on August 14, 2026, against a fifty-two-week high of \$1,998.01. Public market data as of that date. Market data establishes the price levels only; it establishes nothing about why they moved.

³⁸Federal Housing Finance Agency, credit-score policy, as cited above.

I have also applied this framework across an exploratory cross-industry exercise of my own. It reinforced the point above: superficially similar AI-related price movements frequently reflected very different economic mechanisms. I do not offer that exercise as a market-wide empirical result, and nothing in this essay rests on it.

Two boundaries on what the framework itself has established. The first is that this essay proposes a framework rather than claiming to have confirmed it empirically; its investment use is a hypothesis to be tested company by company, and a proper test would state its predictions and its refuting observations before the next observations are collected. The second is narrower. I have no evidence of any company being rerated because AI makes legitimate rights cheaper to investigate or enforce, and I have not tested for it. Minimum Viable Enforcement therefore remains a forward hypothesis.

The framework is not a machine for sorting companies into AI winners and losers; markets already do that passably. Its purpose is to identify which economic variable is actually changing, then ask whether the price carries the right model of its magnitude, its timing, and the long-run economics the valuation assumes.

The opportunity is not to notice that AI matters. It is to understand the resulting economics more accurately than the price does, and to know when you are entitled to act on the difference.

Key takeaway. A correct structural insight is not an investment edge if the price already assumes it.

Ask yourself. What must be true about friction, adaptation, growth and terminal economics for today's price to be right? Which of those assumptions do I actually disagree with?

Part VI. Where Friction Goes Next

When both sides have machines

Once both sides automate, the effects stop being simple.

Effect	What happens
Compression	The cost of overcoming a friction falls sharply
Inversion	A barrier that protected one party starts favoring the other
Creation	New friction appears, often as defensive counter-process
Amplification	An existing right, barrier or process can be exercised repeatedly, or at machine scale
Migration	The bottleneck moves downstream rather than disappearing
Escalation	Both sides invest, one may automate first, and neither gains lasting advantage
Congestion	Cheap submission overwhelms the capacity that processes it

A consumer agent files a claim for \$437. A corporate agent denies it under clause 14. The consumer agent responds that clause 14 does not apply because of a term in the amended schedule. The corporate agent checks and approves. Nine seconds, no humans. That version is genuinely good: small legitimate claims that used to die of attrition now get resolved.

Now change one variable. The corporate agent says the amended schedule was never executed. The consumer agent produces a copy. The corporate agent says the copy is not authentic. The machines have arrived somewhere neither can go, because the question is no longer what the contract means. It is which account of the facts is true, and that has never been settled by arguing more cheaply.

A fall in the cost of making a claim does not imply a fall in the cost of resolving one.

AI can compress micro-friction, the cost of any one transaction, while increasing macro-friction, the load on whatever has to resolve them all. The CFPB's June 2026 statement about its own complaint system is the closest institutional analogy available: submission volume grew to the point that the Bureau said it could no longer rely on the resulting data as a reliable picture of market conditions.³⁹ That episode is not attributed to AI, and I do not attribute it. The inference I draw is narrower and, I think, general: generating assertions can scale faster than institutions can verify and resolve them. Courts, regulators, arbitrators, and auditors were all sized for a slower rate of arrival, and none get faster because the parties in front of them did.

Institutions will not simply absorb this. Cheap challenge invites new verification requirements, identity checks, structured submission formats, rate limits, automated triage and sometimes fees. Some of that will be protective friction doing exactly what protective friction should. Some will become the next asymmetric friction, because a structured-submission standard is easier for an institution to satisfy than for an individual. I state that as a forward inference rather than an observation, but it follows directly from the taxonomy above: compression on one side invites creation on the other. The equilibrium is not friction falling once. It is counterparties repeatedly redesigning where the friction sits.

When intelligence gets cheap, scarcity moves

In the industrial economy the binding scarcity was capital and productive capacity. In the information economy it shifted toward information, expertise, software and distribution. Now a large share of routine cognitive work is getting dramatically cheaper. Not free, and not uniformly: on the one task Dell'Acqua and colleagues engineered to sit beyond the model's capability, AI users were markedly less likely to reach the correct answer, which is a standing reminder that capability has a jagged edge.⁴⁰

As one formerly scarce input gets cheaper, other constraints become relatively more important. So what stays scarce? Ownership and enforceable rights. Trusted relationships and brands that took decades to build. Provenance, meaning credible knowledge of where something came from. Exclusive access, to data, customers, supply or a regulated market. Physical capacity. Distribution. Regulatory permission. And institutional authority: the capacity of a court, a regulator, an auditor or a certifying body to issue a determination other parties will act on.

³⁹CFPB, "The CFPB, as cited above.

⁴⁰Dell'Acqua, F., McFowland III, E., Mollick, E., Lifshitz, H., Kellogg, K. C., as cited above.

If cognition becomes cheap, own what cognition alone cannot manufacture.

Key takeaway. Making claims cheap does not make resolving them cheap. As argument becomes abundant, verification and authority become the scarce inputs.

Ask yourself. If counterparties can generate challenges at machine scale, what becomes your binding verification or settlement capacity? What protective friction will institutions add in response?

Note the last item. **I think AI is commoditizing expertise faster than it is commoditizing authority.** Authority here means the recognized standing to make a determination others will act on. AI can already produce some sophisticated analysis at dramatically lower marginal cost. Producing analysis does not itself confer institutional authority: it cannot produce a judgment, an approval, a certification, or a determination a counterparty is obliged to honor. The first is knowledge. The second is standing.

Which means the constraint moves from *can we produce a good answer* toward *whose answer an institution will stand behind*. Cheap challenge, cheap enforcement, cheap analysis, and cheap counter-analysis all converge on the same bottleneck: not the production of arguments, but the establishment of what should be acted upon. What follows is about who captures value in the meantime.

This is an operating thesis

I said at the start that I am not a disinterested observer.

The doctrine I run companies by has four parts. Destroy the friction we pay, because it is pure cost. Destroy the friction our customers pay that produces no value for them, because someone will and I would rather it be us. Activate the legitimate rights we already own, because rights whose exercise has been uneconomic can contain value the owner has never been able to realize. And accumulate the scarce things cheap intelligence cannot manufacture, because that is where margin is going.

The five tools in Part III and the change-capacity assessment in Part IV are not illustrations. They are what I actually run, against companies I own, companies I am considering, and companies I compete with, and Part V is what happens when I make myself check the answer against price. A framework that can only confirm what its author already believed is not an analytical tool, it is a rhetorical one.

I expect substantial economic value to move as friction is repriced, and I intend to compete for it. I want the companies I build, their shareholders, and my own capital to benefit from understanding this early and acting deliberately. That is a statement of intent and interest, not a prediction of returns. Nothing here is investment advice, a recommendation on any security, or a promise of any outcome, and I hold positions that this framework informs.

The unresolved scarcity

AI will not produce a frictionless economy. That was never the available outcome.

The strategic question is no longer whether AI will reduce your costs. It is which entries on your friction balance sheet are moving, in which direction, whether you can move before your counterparties do, and how much of that is already in the price.

Return to Sarah and David. AI can make Sarah's dispute cheaper to understand, assemble and pursue. It can make David's infringement question cheaper to investigate, and it can do both far better than anything either of them could previously afford. But neither problem ends when the analysis is produced. Sarah still needs an institution to accept that her evidence warrants changing her record. David still needs to decide whether his evidence warrants putting capital, reputation and legal process behind a consequential claim. Their problems begin as friction problems. They end as reliance problems.

I began with a simple observation: AI is making answers cheaper faster than it is making them good enough to act on. Friction is why that matters economically. When search, expertise, monitoring and challenge get cheaper, work that was not worth doing crosses the threshold into action. Dormant rights wake up. Dormant obligations become visible. Competitors move faster, customers scrutinize more, companies enforce more, and eventually machines challenge machines. That does not produce a frictionless economy. It moves friction toward the places where consequences have to become real: evidence, verification, adjudication, authority. Many of the institutions we rely on for those things were designed for a world in which producing sophisticated claims and counterclaims was substantially more expensive, and that assumption is changing quickly.

Which leaves what I think is one of the consequential questions of this period, and the one this essay has been walking toward without answering. What happens when producing an argument, challenging it, and defending against the challenge all become cheap, but establishing what can actually be trusted enough to act upon does not?

That is the problem I have been working on.

About Expound Laboratories



Expound Laboratories develops research, standards, and computational systems concerned with Semantic Computation and Finality Assurance. This article is not about that work. It is about the economics that make the underlying problem worth caring about.

The E^x mark combines the natural exponential with the company's name, and carries change, compounding, and the question of what an expanding quantity of computation actually establishes.

Expound Laboratories, Semantic Computation, Finality Assurance, Continuous Finality, Accepted Work per Dollar and the E^x mark are trademarks of Kenneth Lawrence Cheney, owner of Expound Laboratories and copyright holder of this work. Third-party names are descriptive references to their own public disclosures, with no affiliation or endorsement implied.

This article states the author's economic point of view and operating philosophy. It is not investment advice, not a recommendation with respect to any security, and not a promise of any outcome. The author holds positions that this framework informs.

<https://expound.com> • info@expound.com

Public Article • © 2026 Kenneth L. Cheney. All rights reserved. Marks identified above.